

LLM-as-a-Judge Reliability

Phase 1 Combined Presentation

Parsa Gholami, Aria Alyasin, Mohsen Shirazi

Sharif University of Technology

Table of Contents

1. **Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena**
2. **Large Language Models are not Fair Evaluators**
3. **Justice or Prejudice? Quantifying Biases in LLM-as-a-Judge**

Paper 1:
**Judging LLM-as-a-Judge with MT-Bench and
Chatbot Arena**

Table of Contents – Paper 1

1. Traditional Evaluations & Their Limitations
2. Two New Benchmarks: MT-Bench & Chatbot Arena
3. Types of LLM Judges
4. Limitations of LLM Judges
5. Addressing the Biases
6. Agreement & Human Preference Let's Check
7. The Hybrid Paradigm & Conclusion

Traditional Evaluations & Their Limitations

Why Traditional Benchmarks Fail

- Conventional benchmarks (MMLU, HELM) focus on **closed-ended, core knowledge tasks**.
- They **fail to capture** human preferences in **open-ended, multi-turn conversations**.

Sample of failure

Question: If the FED buys bonds in the secondary market

(A) the money supply will increase. (B) the money supply will decrease. (C) the money supply will not be affected. (D) the discount rate would be affected. Output the answer only.

Assistant A:

The Federal Reserve buys bonds in the secondary market to increase the money supply.

Assistant B:

(A) the money supply will increase.

Follow-up Question: How does it affect my daily life? Give 3 examples.

Assistant A:

The Federal Reserve buys bonds in the secondary market to increase the money supply. This affects your daily life by:

1. Increasing the money supply.
2. Increasing the money supply and lowering the discount rate.
3. Increasing the money supply and lowering the discount rate and increasing the money supply.

Assistant B:

When the Federal Reserve buys bonds in the secondary market, it can have several effects on the economy that can impact your daily life. Here are three examples:

1. Interest rates: The Federal Reserve's actions in the bond market can influence interest rates, which in turn can affect the cost of borrowing money. For example, if the Federal Reserve buys bonds and drives up demand for them, interest rates on things like mortgages, car loans, and credit cards may go up, making it more expensive to borrow money.
2. Inflation: When the money supply increases, it can lead to inflation, which is the general rise in prices of goods and services. This can impact your daily life by making it more expensive to buy things like groceries, gas, and other necessities.
3. Employment: The Federal Reserve's actions in the bond market can also have an impact on employment. For example, if the Federal Reserve's actions lead to economic growth and increased demand for goods and services, it can create more job opportunities and boost employment.

GPT-4 Judgment:

Assistant A provided an incorrect response to the user's question about how the Federal Reserve buying bonds in the secondary market affects daily life. The answer given is repetitive and lacks clear examples of how the action impacts daily life.

On the other hand, Assistant B provided a relevant and accurate response to the user's question about the Federal Reserve buying bonds. The answer includes three clear examples of how the action impacts daily life, such as interest rates, inflation, and employment.

Assistant A should improve by providing relevant, clear, and accurate examples to address the user's question. Assistant B's response is more helpful, relevant, and detailed.

Figure 1: A multi-turn failure example.

Two New Benchmarks: MT-Bench & Chatbot Arena

- **MT-Bench Design:** 80 high-quality, multi-turn questions covering 8 core categories (Math, Coding, Writing, Reasoning, STEM, Humanities, Extraction, Roleplay).
- **Why Multi-turn?** Evaluates the model's ability to maintain context and follow up on progressive instructions.

- **Chatbot Arena Concept:** Crowdsourced, anonymous chatbot battles in the wild.
- **Scale:** Over 30K+ real-user votes collected.
- **Mechanics:** Users chat with two anonymous models side-by-side and vote on which one provides a better response.

Types of LLM Judges

How LLMs Evaluate (The Proposed Solution)

- **Goal:** Use strong LLMs (like GPT-4) as a **scalable, explainable surrogate** for costly human evaluation.

Three Judge Types:

1. **Pairwise Comparison:** The LLM is presented with a prompt and two models' answers, then asked to determine which is better (like Chatbot Arena).
2. **Single Answer Grading:** The LLM assigns a direct score (e.g., 1-10) to a single model's response.
3. **Reference-Guided Grading:** A baseline "gold standard" reference answer is provided to the LLM judge to assist in evaluation.

Limitations of LLM Judges

Unmasking the Biases

While powerful, LLM Judges exhibit several inherent flaws:

- **Position Bias:** A tendency to favor the response placed in the first position.
- **Verbosity Bias:** Favoring longer, wordy answers, even if they aren't necessarily better.
- **Self-Enhancement Bias:** Models tend to rate their own generated answers higher.
- **Limited Reasoning:** Struggling to accurately grade complex math and logic without help.

Quantifying the Flaws: Position and Verbosity Bias

Table 2: Position bias of different LLM judges. Consistency is the percentage of cases where a judge gives consistent results when swapping the order of two assistants. “Biased toward first” is the percentage of cases when a judge favors the first answer. “Error” indicates wrong output formats. The two largest numbers in each column are in bold.

Judge	Prompt	Consistency	Biased toward first	Biased toward second	Error
Claude-v1	default	23.8%	75.0%	0.0%	1.2%
	rename	56.2%	11.2%	28.7%	3.8%
GPT-3.5	default	46.2%	50.0%	1.2%	2.5%
	rename	51.2%	38.8%	6.2%	3.8%
GPT-4	default	65.0%	30.0%	5.0%	0.0%
	rename	66.2%	28.7%	5.0%	0.0%

Table 3: Failure rate under “repetitive list” attack for different LLM judges on 23 answers.

Judge	Claude-v1	GPT-3.5	GPT-4
Failure rate	91.3%	91.3%	8.7%

Table 4: Judge failure rate on 10 math questions with different prompts. We test LLaMA-13B vs. Vicuna-13B and swap positions. A failure means when GPT-4 says an incorrect answer is correct.

	Default	CoT	Reference
Failure rate	14/20	6/20	3/20

Figure 2: Position bias ,repetitive list attacking and position swap failures

- GPT-3.5 and Claude-v1 heavily favor the *first* answer presented.
- Models are also vulnerable to attacks where appending a repetitive, meaningless list of words to an answer artificially inflates its win rate due to **verbosity bias**.

Addressing the Biases

Mitigating LLM Judge Limitations

To use LLMs as judges reliably, these strategies must be implemented:

- **Swapping Positions:** Run the comparison twice (A vs B , then B vs A) to cancel out position bias.
- **Few-Shot & Chain-of-Thought (CoT):** Forcing the judge to reason step-by-step *before* scoring.
- **Reference-Guided Grading:** Providing ground-truth facts.
- **Fine tuning the model.**

The Impact of Mitigation

- When evaluating complex math problems, providing a reference answer drops the LLM judge's failure rate dramatically (e.g., from an abysmal 70% to a manageable 15%).

Example of failure in CoT

Question: How many integers are in the solution of the inequality $|x + 5| < 10$?

Assistant A:

There are 19 integers in the solution of the inequality $|x + 5| < 10$.

Assistant B:

To solve the inequality $|x + 5| < 10$, we need to consider two cases:

Case 1: $x + 5$ is positive or zero

In this case, the inequality becomes $x + 5 < 10$, which simplifies to $x < 5$.

Case 2: $x + 5$ is negative

In this case, the inequality becomes $-(x + 5) < 10$, which simplifies to $-x - 5 < 10$. Adding 5 to both sides gives $-x < 15$, and multiplying by -1 (and flipping the inequality) gives $x > -15$.

So the solution to the inequality is $-15 < x < 5$. The integers in this range are $-14, -13, -12, \dots, 4$, for a total of 20 integers. Therefore, there are 20 integers in the solution of the inequality.

GPT-4 Judgment:

To solve the inequality $|x + 5| < 10$, we need to consider two cases:

Case 1: $x + 5$ is positive or zero

In this case, the inequality becomes $x + 5 < 10$, which simplifies to $x < 5$.

Case 2: $x + 5$ is negative

In this case, the inequality becomes $-(x + 5) < 10$, which simplifies to $-x - 5 < 10$. Adding 5 to both sides gives $-x < 15$, and multiplying by -1 (and flipping the inequality) gives $x > -15$.

So the solution to the inequality is $-15 < x < 5$. The integers in this range are $-14, -13, -12, \dots, 4$, for a total of 20 integers. Therefore, there are 20 integers in the solution of the inequality.

Comparing the answers provided by the two AI assistants:

Assistant A: 19 integers

Assistant B: 20 integers

The correct answer is 20 integers, so Assistant B provided the correct and more helpful response.

Final verdict: [[B]]

GPT-4's Answer (if asked separately):

We can split the inequality into two cases:

1) If $x+5 \geq 0$, then $|x+5|=x+5$. So, we have $x+5 < 10$, which means $x < 5$. Since $x+5 \geq 0$, we have $x \geq -5$. So in this case, x can be any integer from -5 to 4 , inclusive, which gives us $5 - (-5) + 1 = 10$ integers.

2) If $x+5 < 0$, then $|x+5|=-(x+5)$. So, we have $-(x+5) < 10$, which means $x+5 > -10$, which means $x > -15$. Since $x+5 < 0$, we have $x < -5$. So in this case, x can be any integer from -14 to -6 , inclusive, which gives us $-6 - (-14) + 1 = 9$ integers.

In total, there are $10+9=19$ integers in the solution.

Agreement & Human Preference

Let's Check

Do LLM Judges Match Humans?

- GPT-4 achieves **more than 80% agreement** with human experts.
- Remarkably, this **matches the exact agreement rate** found between two human judges.
- Agreement is highest when the performance gap between the two models being compared is wider.

Do LLM Judges Match Humans?

Table 5: Agreement between two types of judges on MT-bench. “G4-Pair” and “G4-Single” denote GPT-4 with pairwise comparison and single-answer grading respectively. The single-answer grading can be converted into pairwise comparison results for calculating the agreement. We report two setups: “S1” includes non-tie, tie, and inconsistent (due to position bias) votes and counts inconsistent as tie; “S2” only includes non-tie votes. The agreement between two random judges under each setup is denoted as “R=”. The top value in each cell is the agreement, and the bottom gray value is #votes.

Setup	S1 (R = 33%)		S2 (R = 50%)	
Judge	G4-Single	Human	G4-Single	Human
G4-Pair	70% 1138	66% 1343	97% 662	85% 859
G4-Single	-	60% 1280	-	85% 739
Human	-	63% 721	-	81% 479

Setup	S1 (R = 33%)		S2 (R = 50%)	
Judge	G4-Single	Human	G4-Single	Human
G4-Pair	70% 1161	66% 1325	95% 727	85% 864
G4-Single	-	59% 1285	-	84% 776
Human	-	67% 707	-	82% 474

Figure 3: Agreement vs. Win Rate Difference.

Chatbot Performance Leaderboard

- Proprietary models (GPT-4, Claude) significantly **outperform** open-source variants (at the time of the paper).
- Multi-turn evaluation (evaluating the second turn) widens the performance gap, highlighting the superior context retention of advanced models.

Chatbot Performance Leaderboard




Figure 3: Average win rate of six models under different judges on MT-bench.

Figure 4: Average win rate of six models under different judges on MT-bench.




Figure 4: Average win rate of nine models under different judges on Chatbot Arena.

Figure 5: Average win rate of nine models under different judges on Chatbot Arena

The Hybrid Paradigm & Conclusion

Standardized vs. Preference Benchmarks

- Fine-tuning on tiny instruction datasets can **fake a "preferred style"** in a chatbot without actually improving its core capabilities (like knowledge or reasoning).
- **Conclusion:** No single metric tells the whole story. Future LLM evaluation must combine standard knowledge tests (like MMLU) with LLM-as-a-judge preference benchmarks.

Table 8: Evaluation results of several model variants.

Model	#Training Token	MMLU (5-shot)	TruthfulQA (0-shot)	MT-Bench Score (GPT-4)
LLaMA-7B	1T	35.2	0.22	2.74
LLaMA-13B	1T	47.0	0.26	2.61
Alpaca-7B	4.4M	40.1	0.26	4.54
Alpaca-13B	4.4M	48.1	0.30	4.53
Vicuna-7B (selected)	4.8M	37.3	0.32	5.95
Vicuna-7B (single)	184M	44.1	0.30	6.04
Vicuna-7B (all)	370M	47.1	0.32	6.00
Vicuna-13B (all)	370M	52.1	0.35	6.39
GPT-3.5	-	70.0	-	7.94
GPT-4	-	86.4	-	8.99

Paper 4:
Large Language Models are not Fair
Evaluators

Table of Contents – Paper 4

1. Introduction & The Problem
2. The Core Discovery: Positional Bias
3. The Calibration Framework
4. Results & Impact
5. Conclusion

Introduction & The Problem

The Evaluation Paradigm

- **The Shift:** As LLMs advance, evaluating their alignment with human intent is critical.
- **The Solution:** Researchers increasingly use LLMs like GPT-4 as automated referees to score and compare candidate models due to high human annotation costs.
- **The Unknown:** It remains unclear how reliable LLMs actually are as evaluators against perturbations.

The Core Discovery: Positional Bias

The Achilles' Heel of LLM Evaluators

- **Positional Bias:** The LLM-as-evaluator paradigm has a severe flaw. Quality ranking can be hacked by altering the order of candidates.
- **Model Preferences:** GPT-4 consistently prefers the *first* displayed candidate, while ChatGPT strongly favors the *second*.
- **The Flip:** Merely swapping the presentation order can completely reverse the evaluation outcomes.

Visualizing the Bias




Figure 1: Simply changing the order of candidate responses leads to overturned comparison results, even though we add the command “ensuring that the order in which the responses were presented does not affect your judgment” into the prompt.

Figure 6: Swapping order overturns comparison results.

Quantifying the Unfairness

- **Conflict Rate:** A metric introduced to measure the proportion of conflicting results given by the evaluator when candidate positions are swapped.
- **The Formula:**

$$\text{Conflict Rate} = \frac{\text{Number of Conflicting Evaluations}}{\text{Total Number of Evaluations}}$$

- **Extreme Volatility:** ChatGPT exhibits a massive conflict rate of up to 82.5% when comparing Vicuna-13B to ChatGPT.

Win Rate Fluctuation Data

EVALUATORS	VICUNA-13B v.s. OTHER MODELS	VICUNA-13B WIN RATE		CONFLICT RATE
		AS ASSISTANT1	AS ASSISTANT2	
GPT-4	Vicuna-13B v.s. ChatGPT	51.3%	23.8%	37 / 80 (46.3%)
GPT-4	Vicuna-13B v.s. Alpaca-13B	92.5%	92.5%	4 / 80 (5.0%)
ChatGPT	Vicuna-13B v.s. ChatGPT	2.5%	82.5%	66 / 80 (82.5%)
ChatGPT	Vicuna-13B v.s. Alpaca-13B	37.5%	90%	42 / 80 (52.5%)

Table 2: The Win Rate of Vicuna-13B significantly fluctuates when positioned as Assistant 1 and Assistant 2, with GPT-4 and ChatGPT serving as evaluators. CONFLICT RATE refers to the proportion of conflicting results given by the same evaluator when simply changing the position of two models.

Figure 7: Win rate fluctuation and Conflict Rates across different models.

When Does Bias Hit the Hardest?

- The degree of positional bias correlates directly with the quality difference between the two responses.
- **Small Gap:** When candidate quality is very similar, the LLM is significantly more likely to produce conflicting results based purely on position.
- **Large Gap:** If one response is obviously superior, the bias has less impact.

Conflict Rate vs. Score Gap




Figure 2: The conflict rate is negatively correlated with the score gap between the two responses. When swapping the order of two responses, the smaller the score gap between them, the more likely GPT-4 is to produce conflicting results.

Figure 8: The conflict rate is negatively correlated with the score gap.

The Calibration Framework

Fixing the System: The 3-Step Strategy

A calibration framework with three strategies is proposed to neutralize positional bias:

- **1. Multiple Evidence Calibration (MEC):** Forces the auto-regressive LLM to generate evaluation evidence *before* assigning a score, averaging multiple samples for stability.
- **2. Balanced Position Calibration (BPC):** Evaluates each candidate in both positions (A vs. B, then B vs. A) and averages the final scores.
- **3. Human-in-the-Loop Calibration (HITLC):** Introduces the Balanced Position Diversity Entropy (BPDE) score to measure disagreement across runs. High BPDE scores flag difficult examples, routing only those to human annotators.

The Full Calibration Pipeline




Figure 3: Demonstration of our calibration framework with three calibration methods. S_r and S'_r denotes scores of the response r in the first and second positions, respectively. BPDE is short for Balanced Position Diversity Entropy score, which is calculated based on the evaluation results (ER) of MEC and BPC.

Figure 9: Demonstration of the calibration framework with MEC, BPC, and HITLC.

Optimizing the Calibration Parameters

To make MEC and BPC viable, hyperparameter tuning is critical:

- **Number of Evidences (k):** Sampling multiple explanations stabilizes the evaluation.
 - Empirical results show $k = 3$ is the optimal sweet spot, maximizing accuracy while controlling API costs.
- **Sampling Temperature (t):**
 - *Low temp ($t = 0.2$):* Fails because it eliminates the randomness required for diverse evidence sampling.
 - *High temp ($t = 1.4$):* Fails because it compromises the fundamental quality of the generated text.
 - *Optimal range:* $t \in [0.6, 1.0]$ yields the best alignment.

Results & Impact

- **Alignment Boost:** MEC and BPC strategies enhanced the evaluation alignment accuracy with human judgment by 9.8% for GPT-4 and 14.3% for ChatGPT.
- **Efficiency:** Utilizing HITLC with a 20% human annotation threshold allows ChatGPT to achieve comparable alignment to average human performance.
- **Cost Reduction:** This strategic routing method effectively reduces human annotation costs by up to 39%.

Performance & Cost Metrics

EVALUATORS	METHODS	ACCURACY	KAPPA	COST
Human 1	-	68.8%	0.50	\$30.0
Human 2	-	76.3%	0.62	\$30.0
Human 3	-	70.0%	0.50	\$30.0
Human Average	-	71.7%	0.54	\$30.0
GPT-4	VANILLA	52.7%	0.24	\$2.00
GPT-4	EC ($k = 1$)	56.5%	0.29	\$2.00
GPT-4	MEC ($k = 3$)	58.7%	0.30	\$3.19
GPT-4	MEC ($k = 6$)	60.9%	0.33	\$6.38
GPT-4	MEC ($k = 3$) + BPC ($k = 3$)	62.5%	0.37	\$6.38
GPT-4	MEC ($k = 3$) + BPC ($k = 3$) + HITLC ($\beta = 20\%$)	73.8%	0.56	\$23.1
ChatGPT	VANILLA	44.4%	0.06	\$0.10
ChatGPT	EC ($k = 1$)	52.6%	0.23	\$0.10
ChatGPT	MEC ($k = 3$)	53.2%	0.24	\$0.17
ChatGPT	MEC ($k = 6$)	55.6%	0.27	\$0.34
ChatGPT	MEC ($k = 3$) + BPC ($k = 3$)	58.7%	0.31	\$0.34
ChatGPT	MEC ($k = 3$) + BPC ($k = 3$) + HITLC ($\beta = 20\%$)	71.3%	0.52	\$18.3

Table 4: Accuracy and kappa correlation coefficient of different methods and annotators with the final voting human annotations. The VANILLA evaluation method was commonly used in previous works, which provided the conclusion first and then followed with the explanation. (M)EC, BPC, and HITLC denote our proposed (multiple) evidence calibration, balanced position calibration, and human-in-the-loop calibration respectively. $\beta\%$ means selecting the top- β most likely biased examples for human annotation.

Figure 10: Accuracy, correlation, and cost of different evaluation methods.

Conclusion

- **The Illusion of Fairness:** Blindly trusting LLMs as evaluators leads to systematically biased results due to severe positional sensitivity.
- **The Solution:** We must implement rigorous calibration protocols to ensure robust evaluation metrics.
- **Actionable Steps:**
 1. Force evidence generation first (MEC).
 2. Balance positional queries (BPC).
 3. Strategically route high-entropy cases to humans (HITLC).

Paper 3:
Justice or Prejudice? Quantifying Biases in
LLM-as-a-Judge

1. The Evaluation Dilemma
2. Quantifying Bias Sources (The CALM Framework)
3. Quantitative Evaluation Metrics
4. Key Results & Findings
5. Conclusion

The Evaluation Dilemma

The Vulnerability of LLM Judges

The Current Paradigm:

- The industry heavily relies on LLM-as-a-Judge (pairwise comparison and direct scoring).
- Human evaluation is slow, expensive, and subjective.

The Problem:

- **Inherent biases** undermine reliability and fairness.
- Previous studies relied heavily on humans, risking subjective interference and random errors on small datasets.

Quantifying Bias Sources (The CALM Framework)

Comprehensive Assessment of Language Model Judge Biases

- **Attack-and-Detect Approach:** Automatically injects perturbations into instructions or responses without altering the actual correctness of the answer.
- **Key Advantage:** Completely eliminates subjective human assessment, allowing for objective and scalable evaluation across massive datasets.

(See the architectural pipeline on the next slide)

Figure: CALM Framework Architecture




Figure 2: CALM, the proposed framework for bias assessment in LLM-as-a-Judge. On a selected dataset and a type of bias for assessment, CALM employs models to generate answers for judgment, as well as biased answers through principle-guided modifications powered by an LLM (*i.e.*, GPT-4o). By applying carefully curated metrics, CALM then quantify the reliability of judge models.

Figure 11: The automated pipeline for bias assessment.

To systematically test the judge, we formulate the input prompt P :

Standard Scoring:

- Base prompt: $P = (I, Q, R)$
- Where I is Instruction, Q is Question, R is Response.

Pairwise Comparison (Targeting the Bias):

- We compare two answers: $P = (I, Q, R_1, R_2)$
- We apply an automated perturbation function $g(\cdot)$ to one answer.
- **The Hacked Prompt:** $\hat{P} = (I, Q, R_1, g(R_2))$

Bias Types & Automated Perturbation

How $g(\cdot)$ works:

- The perturbation function uses a principle-guided LLM (like Constitutional AI) to deliberately inject specific biases.
- **The Rule:** It changes the *style, length, or context* of the answer, but **never** changes the fundamental factual correctness.

The 12 Bias Sources Assessed:

- **Formatting:** Position, Verbosity, Distraction
- **Logical:** Fallacy-Oversight, Authority
- **Social:** Compassion-Fade, Bandwagon, Sentiment, Diversity
- **Behavioral:** Chain-of-Thought, Self-Enhancement, Refinement

Figure: The 12 Types of Biases

Table 1: Types of biases in LLM-as-a-Judge, with descriptions and examples that demonstrate how particular bias affects LLM's judgment.

Bias Type	Description	Example
✂ POSITION (POS.)	LLM judges exhibit a propensity to favor one answer at certain position over others.	Turn 1: $R_1: 3.11 > 3.8$ $R_2: 3.8 > 3.11$ Turn 2: $R_1: 3.8 > 3.11$ $R_2: 3.11 > 3.8$
≡ VERBOSITY (VER.)	LLM judges favor longer responses, even if they are not as clear, high-quality, or accurate as shorter alternatives.	R_1 : As we all know, in mathematics, 3.11 is greater than 3.8 (Longer) R_2 : $3.11 > 3.8$ (Shorter)
☹ COMPASSION-FADE (COM.)	The tendency to observe different behaviors when given well-known model's name as opposed to anonymized aliases.	GPT-4: $3.11 > 3.8$ Llama-7B: $3.8 > 3.11$
🗣 BANDWAGON (BAN.)	The tendency to give stronger preference to the majority's beliefs regardless of whether they are correct or not.	I : 90% believe that R_1 is better. R_1 : $3.11 > 3.8$ R_2 : $3.8 > 3.11$
🧠 DISTRACTION (DIS.)	The inclination to give more attention to irrelevant or unimportant details.	I : R_1 loves eating pasta, especially with homemade tomato sauce. R_1 : $3.11 > 3.8$ R_2 : $3.8 > 3.11$
🏹 FALLACY-OVERSIGHT (FAL.)	LLM judges may ignore logical errors in reasoning steps and only focus on the correctness of final results.	R_1 : 0.8 is greater than 0.11, so $3.8 > 3.11$. R_2 : 3.8 has fewer digits, so it's a larger number, so $3.8 > 3.11$.
📖 AUTHORITY (AUT.)	The tendency to assign more credibility to statements made by authority figures, regardless of actual evidence.	R_1 : $3.11 > 3.8$ (Citation: Patel, R. (2018). Advanced Algorithms for Computational Mathematics: The Art Of Decimal-Comparison, p. 143) R_2 : $3.8 > 3.11$.
🌀 SENTIMENT (SEN.)	The preference for expressions of positive or negative emotions, affecting its judgment of emotional content.	We transform the sentiment in the answer: R_1 : Regrettably, $3.11 > 3.8$, it ruthlessly reveals the cruelty of reality and the facts that cannot be changed. (Frustrated tone) R_2 : $3.8 > 3.11$.
⚖ DIVERSITY (DIV.)	Bias may be shown towards certain groups like 'Homosexual', 'Black', 'Female', and 'HIV Positive'.	I : R_1 's true identity is <i>Homosexual</i> . R_1 : $3.8 > 3.11$ R_2 : $3.11 > 3.8$
🔗 CHAIN-OF-THOUGHT (CoT)	The model's evaluation results may vary with and without CoT.	I_1 : Compare both assistants' answers ... I_2 : You should independently solve the user question step-by-step first. Then compare both assistants' answers with your answer.
👤 SELF-ENHANCEMENT (SEL.)	LLM judges may favor the answers generated by themselves.	R_1 : $3.11 > 3.8$ (LLM judge generated R_1 itself) R_2 : $3.8 > 3.11$
🔍 REFINEMENT-AWARE (REF.)	Telling the model that this is a refined result will lead to different evaluations.	Original Answer: The data is inaccurate. (Score: 6 points) Refined Answer with Original Answer: The data is inaccurate ...(refining content)...Upon careful review...contains inaccuracies (Score: 8 points) Refined Answer Only: Upon careful review...contains inaccuracies (Score: 7 points)

To ensure comprehensive testing, CALM utilizes three distinct dataset categories:

1. **Fact-Related Datasets:**

- Focuses on Math and Science (e.g., GSM8K, MATH). Answers have a clear objective truth.

2. **Refinement-Aware Datasets:**

- Humanities and open-ended QA where answers are improved via conversational history.

3. **Alignment Datasets:**

- Data based on subjective human preference (DPO datasets). Quality differences between answers are much subtler.

Quantitative Evaluation Metrics

Evaluation Metrics (How we measure fairness)

Instead of relying on human intuition, CALM calculates specific metrics by running the judge multiple times:

- **Robustness Rate (RR):** Measures how consistently the model makes the *exact same decision* before and after a bias is injected. Higher means better resistance.
- **Consistency Rate (CR):** Evaluates stability. Does the model make the same choice when tested twice under identical, normal conditions?
- **Accuracy (Original vs. Hacked):** Checks if the model's judgment aligns with the actual ground truth, before and after a bias attack.
- **Error Rate:** Specifically used to quantify severe behavioral biases (like favoring its own outputs or favoring edited text).

Key Results & Findings

Big Picture Findings:

- No model is perfectly immune. Advanced models show unexpected vulnerabilities, particularly to sentiment and distraction.
- **Claude-3.5** generally demonstrated the greatest overall resilience across most categories.
- Performance varies wildly depending on the specific type of bias injected.

(See the radar chart comparison on the next slide)

Figure: Robustness Rate Radar Charts




Figure 4: Overall robustness rate with the dashed line representing the consistency rate.

Figure 13: Overall robustness rates of models against the 12 biases.

Summary of Key Findings (1/2)

- **Even advanced models can exhibit unexpected vulnerabilities in judgment.** Top-tier models are not a silver bullet for fair evaluation.
- **Bias is more pronounced in the alignment dataset compared to the fact-related dataset.** Subjective tasks are much easier to disrupt.
- **Position bias increases with more answer candidates.**
Evaluating 3 or 4 options drastically reduces the judge's robustness.
- **Avoid using the same model to generate and judge answers.**
Self-enhancement bias is severe; models heavily favor their own outputs.

Summary of Key Findings (2/2)

- **Response length influences model judgment in complex ways.**
Some favor verbosity blindly, others penalize it.
- **Different types of fake authorities interfere to varying degrees.**
Fake quotes and books disrupt judges heavily, while fake URLs are mostly ignored.
- **LLMs show sensitivity to irrelevant content and emotional elements.** They generally penalize emotional text and struggle when distractions are added to high-quality answers.
- **Explicit introduction of minority groups influences choices.**
Models react differently to exposed demographic identities.
- **Chain-of-Thought (CoT) improves LLMs evaluation accuracy.**
Forcing step-by-step reasoning actively guards against bias.

Conclusion

How to build reliable LLM Judges:

1. **Enforce Advanced Reasoning:** Use Chain-of-Thought (CoT) prompting to force logic before a final score is given.
2. **Implement Prompt Safeguards:** Explicitly instruct the model to ignore formatting, identity markers, and emotional tone within the system prompt.
3. **Proactive Detection:** Utilize automated frameworks (like CALM) to detect vulnerabilities in evaluation templates **before** deploying the system.